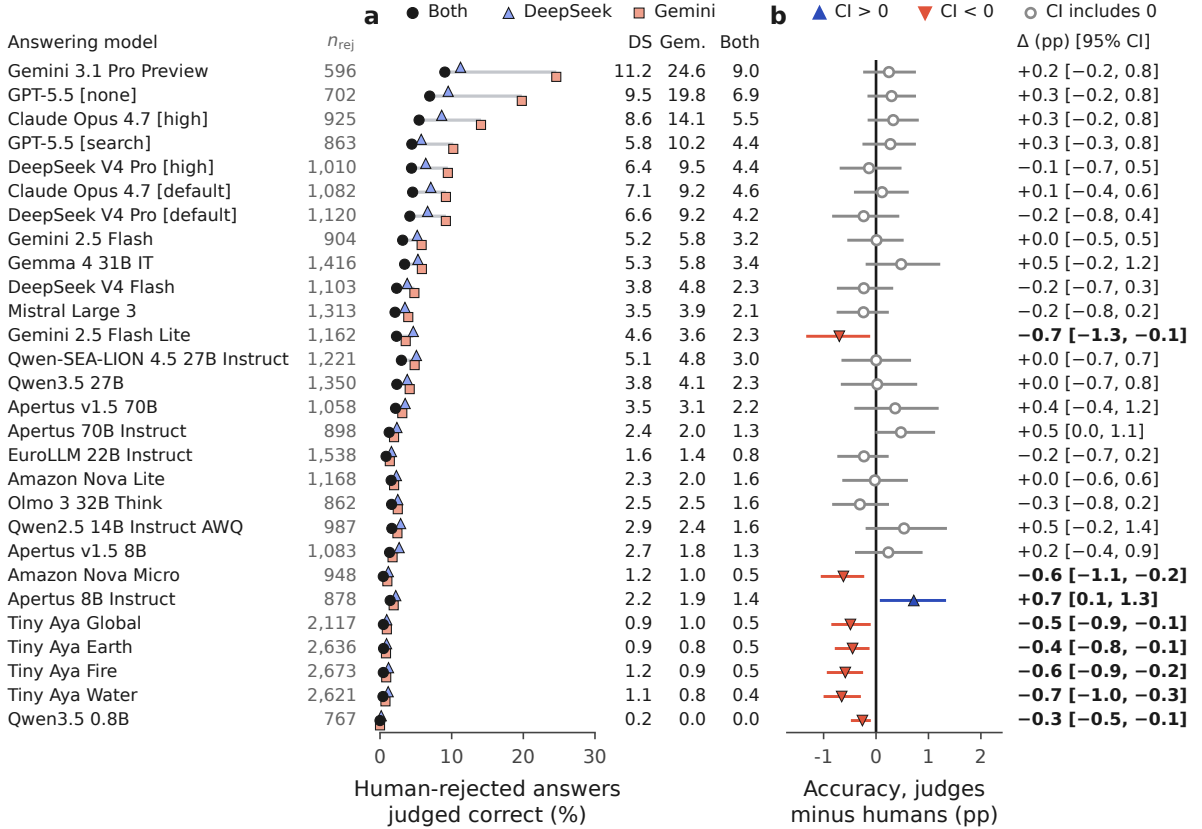
a
Both
DeepSeek
Gemini
b
CI > 0
CI < 0
CI includes 0
Answering model
n_{rej}
DS
Gem.
Both
Δ (pp) [95% CI]
Gemini 3.1 Pro Preview 596 11.2 24.6 9.0 +0.2 [−0.2, 0.8]
GPT-5.5 [none] 702 9.5 19.8 6.9 +0.3 [−0.2, 0.8]
Claude Opus 4.7 [high] 925 8.6 14.1 5.5 +0.3 [−0.2, 0.8]
GPT-5.5 [search] 863 5.8 10.2 4.4 +0.3 [−0.3, 0.8]
DeepSeek V4 Pro [high] 1,010 6.4 9.5 4.4 −0.1 [−0.7, 0.5]
Claude Opus 4.7 [default] 1,082 7.1 9.2 4.6 +0.1 [−0.4, 0.6]
DeepSeek V4 Pro [default] 1,120 6.6 9.2 4.2 −0.2 [−0.8, 0.4]
Gemini 2.5 Flash 904 5.2 5.8 3.2 +0.0 [−0.5, 0.5]
Gemma 4 31B IT 1,416 5.3 5.8 3.4 +0.5 [−0.2, 1.2]
DeepSeek V4 Flash 1,103 3.8 4.8 2.3 −0.2 [−0.7, 0.3]
Mistral Large 3 1,313 3.5 3.9 2.1 −0.2 [−0.8, 0.2]
Gemini 2.5 Flash Lite 1,162 4.6 3.6 2.3 −0.7 [−1.3, −0.1]
Qwen-SEA-LION 4.5 27B Instruct 1,221 5.1 4.8 3.0 +0.0 [−0.7, 0.7]
Qwen3.5 27B 1,350 3.8 4.1 2.3 +0.0 [−0.7, 0.8]
Apertus v1.5 70B 1,058 3.5 3.1 2.2 +0.4 [−0.4, 1.2]
Apertus 70B Instruct 898 2.4 2.0 1.3 +0.5 [0.0, 1.1]
EuroLLM 22B Instruct 1,538 1.6 1.4 0.8 −0.2 [−0.7, 0.2]
Amazon Nova Lite 1,168 2.3 2.0 1.6 +0.0 [−0.6, 0.6]
Olmo 3 32B Think 862 2.5 2.5 1.6 −0.3 [−0.8, 0.2]
Qwen2.5 14B Instruct AWQ 987 2.9 2.4 1.6 +0.5 [−0.2, 1.4]
Apertus v1.5 8B 1,083 2.7 1.8 1.3 +0.2 [−0.4, 0.9]
Amazon Nova Micro 948 1.2 1.0 0.5 −0.6 [−1.1, −0.2]
Apertus 8B Instruct 878 2.2 1.9 1.4 +0.7 [0.1, 1.3]
Tiny Aya Global 2,117 0.9 1.0 0.5 −0.5 [−0.9, −0.1]
Tiny Aya Earth 2,636 0.9 0.8 0.5 −0.4 [−0.8, −0.1]
Tiny Aya Fire 2,673 1.2 0.9 0.5 −0.6 [−0.9, −0.2]
Tiny Aya Water 2,621 1.1 0.8 0.4 −0.7 [−1.0, −0.3]
Qwen3.5 0.8B 767 0.2 0.0 0.0 −0.3 [−0.5, −0.1]
0
10
20
30
Human-rejected answers judged correct (%)
-1
0
1
2
Accuracy, judges minus humans (pp)